



PennState

CMPSC 311 - Introduction to Systems Programming

Assignment Project Exam Help

<https://powcoder.com>

Introduction to Profiling

Add WeChat powcoder

Suman Saha

(Slides are mostly by Professors Patrick McDaniel and
Abutalib Aghayev)



Program Performance



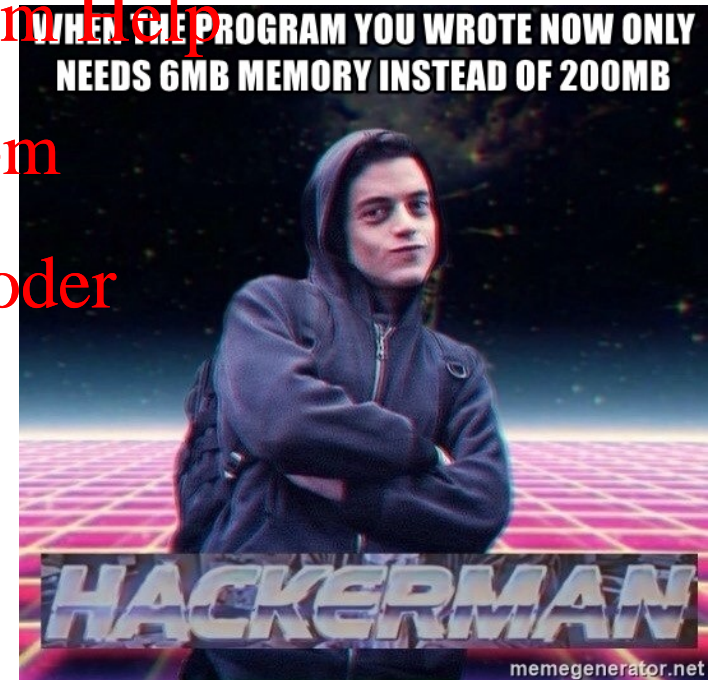
PennState

- Programs run only as well as the code you write
- Poor code often runs poorly
 - Crashes or generates incorrect output (bugs)
 - Is laggy, jittery or slow (inefficient code)
 - Too slow on real inputs (data processing)
 - Not-reactive enough to be usable (interfaces)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Optimization



PennState

- Optimization is the process where you take an existing program and alter it to remove inefficiencies.

- Change algorithms
- Restructure code
- Redesign data structures
- Refactor code

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Career Notes:

1. Learning to optimize your code is essential to becoming a professional programmer.
2. Optimizing code is a phase of development you don't experience in school.
3. You will spend a good deal of your professional life optimizing existing code without changing its function.

Don't optimize prematurely



PennState


- The first rule of optimization is:
 - Don't do it.
- The second rule of optimization (for experts only):
 - Don't do it yet. Measure twice, optimize once.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- It is far, far easier to make a correct program fast than it is to make a fast program correct




Premature optimization is the root of all..



PennState

- "Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered. We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil (quoting Tony Hoare). Yet we should not pass up our opportunities in that critical 3%."



Donald Knuth

Example Inefficient Code



PennState

```
uint64_t prod_one(uint64_t a, uint64_t b) {
    uint64_t out = 0;
    for (uint64_t i = 0; i < a; ++i)
        out += b;
    return out;
}

uint64_t prod_two(uint64_t a, uint64_t b) {
    return a * b;
}

int main(void) {
    uint64_t a = 10000000000, b = 100000000, sum = 0;
    clock_t start;
    double elapsed;

    start = clock();
    sum = prod_one(a, b);
    elapsed = (clock() - start) / CLOCKS_PER_SEC;
    printf("%ld * %ld = %ld (took %.3f seconds) \n", a, b, sum, elapsed);

    start = clock();
    sum = prod_two(a, b);
    elapsed = (clock() - start) / CLOCKS_PER_SEC;
    printf("%ld * %ld = %ld (took %.3f seconds) \n", a, b, sum, elapsed);

    return 0;
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Try it with
 - gcc -O0
 - gcc -O1
 - gcc -O2

```
$ gcc -O0 -Wall opt.c -o opt
[0s euclid:~/tmp (master)]
$ ./opt
10000000000 + 100000000 = 1000000000000000000 (took 25.000 seconds)
10000000000 + 100000000 = 1000000000000000000 (took 0.000 seconds)
[26s euclid:~/tmp (master)]
```

```
$ gcc -O1 -Wall opt.c -o opt
[0s euclid:~/tmp (master)]
$ ./opt
10000000000 + 100000000 = 1000000000000000000 (took 3.000 seconds)
10000000000 + 100000000 = 1000000000000000000 (took 0.000 seconds)
[3s euclid:~/tmp (master)]
```

```
$ gcc -O2 -Wall opt.c -o opt
[0s euclid:~/tmp (master)]
$ ./opt
10000000000 + 100000000 = 1000000000000000000 (took 0.000 seconds)
10000000000 + 100000000 = 1000000000000000000 (took 0.000 seconds)
[0s euclid:~/tmp (master)]
```

Profiling



- Debugging helps the programmer find and fix bugs ...
- **Profiling** helps the programmer find inefficiencies
 - Profiling involves running a version of the program instrumented with code to measure how much time is spent in certain areas of the code.
 - How much time in each of the modules of the program?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

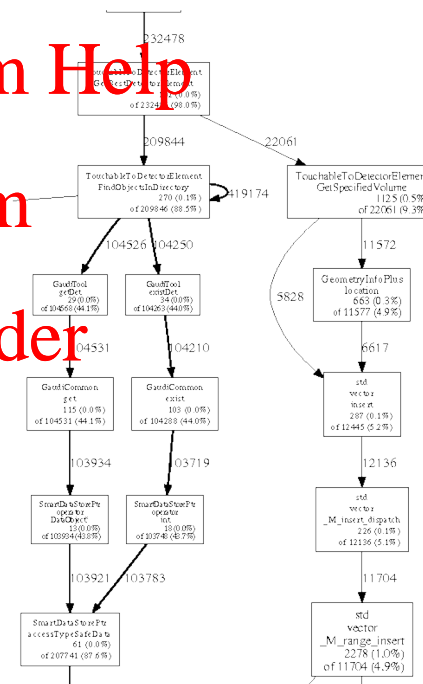


gprof



PennState

- **gprof** is a utility that measures a program's performance and behavior.
- This produces **non-interactive** profile output
- The output provides detail
 - The time that the program ran, and time for each function
 - Statistics and detail on "performance"
 - Which functions called other function and how many times
 - Statistics and detail on the "call graph"



Running gprof



PennState

1. First compile program using the “-pg” flag:

```
$ gcc -pg profiling.c -o profiling
```

2. Run the program (will generate file gmon.out):

```
$ ./profiling
```

3. Run gprof with the named program

```
$ gprof profiling | less
```

4. Review the output

```
$ gprof profiling
Flat profile:
...
```

5. Optimize the program, re-profile
6. GOTO step#1

Gprof (flat profile)



PennState

Assignment Project Exam Help

```
$ gprof opt
Flat profile:
```

```
Each sample counts as 0.01 seconds.
```

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
101.31	21.87	21.87	1	21.87	21.87	prod_one
0.00	21.87	0.00	1	0.00	0.00	prod_two

%
time the percentage of the total running time of the
 program used by this function.

cumulative a running sum of the number of seconds accounted
seconds for by this function and those listed above it.

self
seconds the number of seconds accounted for by this
 function alone. This is the major sort for this
 listing.

Gprof (Call Graph)



PennState

Assignment Project Exam Help

Call graph (explanation follows)

granularity: each sample hit covers 2 byte(s) for 0.05% of 21.87 seconds

index	% time	self	children	called	name
		21.87	0.00	1/1	main [2]
[1]	100.0	21.87	0.00	1	prod_one [1]

[2]	100.0	0.00	21.87		<spontaneous>
		21.87	0.00	1/1	main [2]
				1/1	prod_one [1]
		0.00	0.00	1/1	prod_two [3]

		0.00	0.00	1/1	main [2]
[3]	0.0	0.00	0.00	1	prod_two [3]


Optimization revisited ...



PennState

- When optimizing, you focus on modules of the program which implement the features and processing of the program.

- Which parts of the program you select depends on what parts are running the most.
- then, focus on those parts which take up the most time.



- Profiling tells us where to spend our time.

Amdahl's Law



PennState

- Amdahl's law models the maximum performance gain that can be expected by improving part of the system, i.e., what can we expect in terms of improvement.

Assignment Project Exam Help

- Consider

- k is the percentage of total execution time spent in the optimized module(s).
- s is the execution time expressed in terms of a n-factor speedup (2X, 3X...), which can be found as

<https://powcoder.com>
Add WeChat powcoder

$$s = \frac{\text{original execution time}}{\text{improved execution time}}$$

Amdahl's Law (cont.)



PennState

- The overall speedup T of the program is expressed :

$$T = \frac{1}{(1 - k) + \frac{k}{s}}$$

- Intuition:

<https://powcoder.com>

- $1 - k$ is the part of the program that unchanged

Add WeChat powcoder

- $\frac{k}{s}$ is the ratio of altered program size to speedup

Amdahl's Law (example)



- Assume that a module A of a program is optimized.
 - A represents 45% of the run time of the program.
 - The optimization reduces the runtime of module from 750ms to 50ms.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- What is the program speedup?

Amdahl's Law (example)



- Assume that a module A of a program is optimized.

- A represents 45% of the run time of the program.
- The optimization reduces the runtime of module from 750ms to 50ms.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

$$T = \frac{1}{(1-.45) + \frac{.45}{15}} = \frac{1}{.55 + .03} = 1.724$$

- What is the program speedup? (A: 1.724X)

A more complex example



PennState

- Memory operations currently take 30% of execution time.
- A new widget called a “cache” speeds up 80% of memory operations by a factor of 4.
- A second new widget called a “L2 cache” speeds up 1/2 the remaining 20% by a factor of 2.
- What is the total speed up?

Multiple optimizations: The right way

- We can apply the law for multiple optimizations
- Optimization 1 speeds up x_1 of the program by S_1
- Optimization 2 speeds up x_2 of the program by S_2

$$S_{\text{tot}} = 1/(x_1/S_1 + x_2/S_2 + (1-x_1-x_2))$$

Note that x_1 and x_2 must be disjoint! i.e., S_1 and S_2 must not apply to the same portion of execution.

- Combine both the L1 and the L2
 - memory operations = 0.3
 - $S_{L1} = 4$
 - $x_{L1} = 0.3 * 0.8 = .24$
 - $S_{L2} = 2$
 - $x_{L2} = 0.3 * (1 - 0.8) / 2 = 0.03$
 - $S_{\text{totL2}} = 1/(x_{L1}/S_{L1} + x_{L2}/S_{L2} + (1 - x_{L1} - x_{L2}))$
 - $S_{\text{totL2}} = 1/(0.24/4 + 0.03/2 + (1 - .24 - 0.03))$
 $= 1/(0.06 + 0.015 + .73) = 1.24 \text{ times}$

A even more complex example:



PennState

- Assume another system: we have 4 modules each being measured before and after optimization

Module	Before Optimization (usec)	After Optimization (usec)
A	200	60
B	450	11
C	1000	600
D	125	1

- Now suppose that the runtime of the original execution is 2000 usec, what is the speedup?

$$T = \frac{1}{(1 - k) + \frac{k}{s}}$$

What is going on?



PennState

```
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
```

```
#define SIZE 35000
```

```
int main(void) {
    uint64_t **a = malloc(SIZE * sizeof(uint64_t *));
    for (int i = 0; i < SIZE; ++i)
        a[i] = malloc(SIZE * sizeof(uint64_t));

    uint64_t sum = 0;
    for (int i = 0; i < SIZE; ++i)
        for (int j = 0; j < SIZE; ++j)
            sum += a[i][j];

    printf("sum is %ld\n", sum);

    return 0;
}
```

```
$ gcc -Wall -O2 c.c -o c
$ time ./c
sum is 0
```

```
real    0m1.560s
user    0m1.060s
sys     0m0.500s
```

```
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
```

```
#define SIZE 35000
```

```
int main(void) {
    uint64_t **a = malloc(SIZE * sizeof(uint64_t *));
    for (int i = 0; i < SIZE; ++i)
        a[i] = malloc(SIZE * sizeof(uint64_t));

    uint64_t sum = 0;
    for (int i = 0; i < SIZE; ++i)
        for (int j = 0; j < SIZE; ++j)
            sum += a[j][i];

    printf("sum is %ld\n", sum);

    return 0;
}
```

```
$ gcc -Wall -O2 c.c -o c
$ time ./c
sum is 0
```

```
real    0m14.543s
user    0m13.597s
sys     0m0.948s
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder